

# **COMP2310/COMP6310**

## **Systems, Networks, & Concurrency**


Convener: Shoaib Akram

# Outline

- **Getting content on the web: Telnet/cURL Demo**
  - How the web really works
- Networking Basics
- Proxy
  - Due Tuesday, December 8<sup>th</sup>
  - Grace days allowed
- String Manipulation in C

# The Web in a Textbook

- Client request page, server provides, transaction done.



- A sequential server can handle this. We just need to serve one page at a time.
- This works great for simple text pages with embedded styles.

# Telnet/Curl Demo

## ■ Telnet

- Interactive remote shell – like ssh without security
- Must build HTTP request manually
  - This can be useful if you want to test response to malformed headers

```
[rjaganna@makoshark ~]% telnet www.cmu.edu 80
```

```
Trying 128.2.42.52...
```

```
Connected to WWW-CMU-PROD-VIP.ANDREW.cmu.edu (128.2.42.52).
```

```
Escape character is '^]'.
```

```
GET http://www.cmu.edu/ HTTP/1.0
```

```
HTTP/1.1 301 Moved Permanently
```

```
Date: Sat, 11 Apr 2015 06:54:39 GMT
```

```
Server: Apache/1.3.42 (Unix) mod_gzip/1.3.26.1a mod_pubcookie/3.3.4a mod_ssl/2.8.31 OpenSSL/0.9.8e-fips-rhel5
```

```
Location: http://www.cmu.edu/index.shtml
```

```
Connection: close
```

```
Content-Type: text/html; charset=iso-8859-1
```

```
<!DOCTYPE HTML PUBLIC "-//IETF//DTD HTML 2.0//EN">
```

```
<HTML><HEAD>
```

```
<TITLE>301 Moved Permanently</TITLE>
```

```
</HEAD><BODY>
```

```
<H1>Moved Permanently</H1>
```

```
The document has moved <A HREF="http://www.cmu.edu/index.shtml">here</A>.<P>
```

```
<HR>
```

```
<ADDRESS>Apache/1.3.42 Server at <A HREF="mailto:webmaster@andrew.cmu.edu">www.cmu.edu</A> Port 80</ADDRESS>
```

```
</BODY></HTML>
```

```
Connection closed by foreign host.
```

# Telnet/cURL Demo

## ■ cURL

- “URL transfer library” with a command line program
- Builds valid HTTP requests for you!

```
[rjaganna@makoshark ~]% curl http://www.cmu.edu/
<!DOCTYPE HTML PUBLIC "-//IETF//DTD HTML 2.0//EN">
<HTML><HEAD>
<TITLE>301 Moved Permanently</TITLE>
</HEAD><BODY>
<H1>Moved Permanently</H1>
The document has moved <A HREF="http://www.cmu.edu/index.shtml">here</A>.<P>
<HR>
<ADDRESS>Apache/1.3.42 Server at <A HREF="mailto:webmaster@andrew.cmu.edu">www.cmu.edu</A> Port
80</ADDRESS>
</BODY></HTML>
```

- Can also be used to generate HTTP proxy requests:

```
[rjaganna@makoshark ~]% curl --proxy lemonshark.ics.cs.cmu.edu:3092 http://www.cmu.edu/
<!DOCTYPE HTML PUBLIC "-//IETF//DTD HTML 2.0//EN">
<HTML><HEAD>
<TITLE>301 Moved Permanently</TITLE>
</HEAD><BODY>
<H1>Moved Permanently</H1>
The document has moved <A HREF="http://www.cmu.edu/index.shtml">here</A>.<P>
<HR>
<ADDRESS>Apache/1.3.42 Server at <A HREF="mailto:webmaster@andrew.cmu.edu">www.cmu.edu</A> Port
80</ADDRESS>
</BODY></HTML>
```

# How the Web Really Works

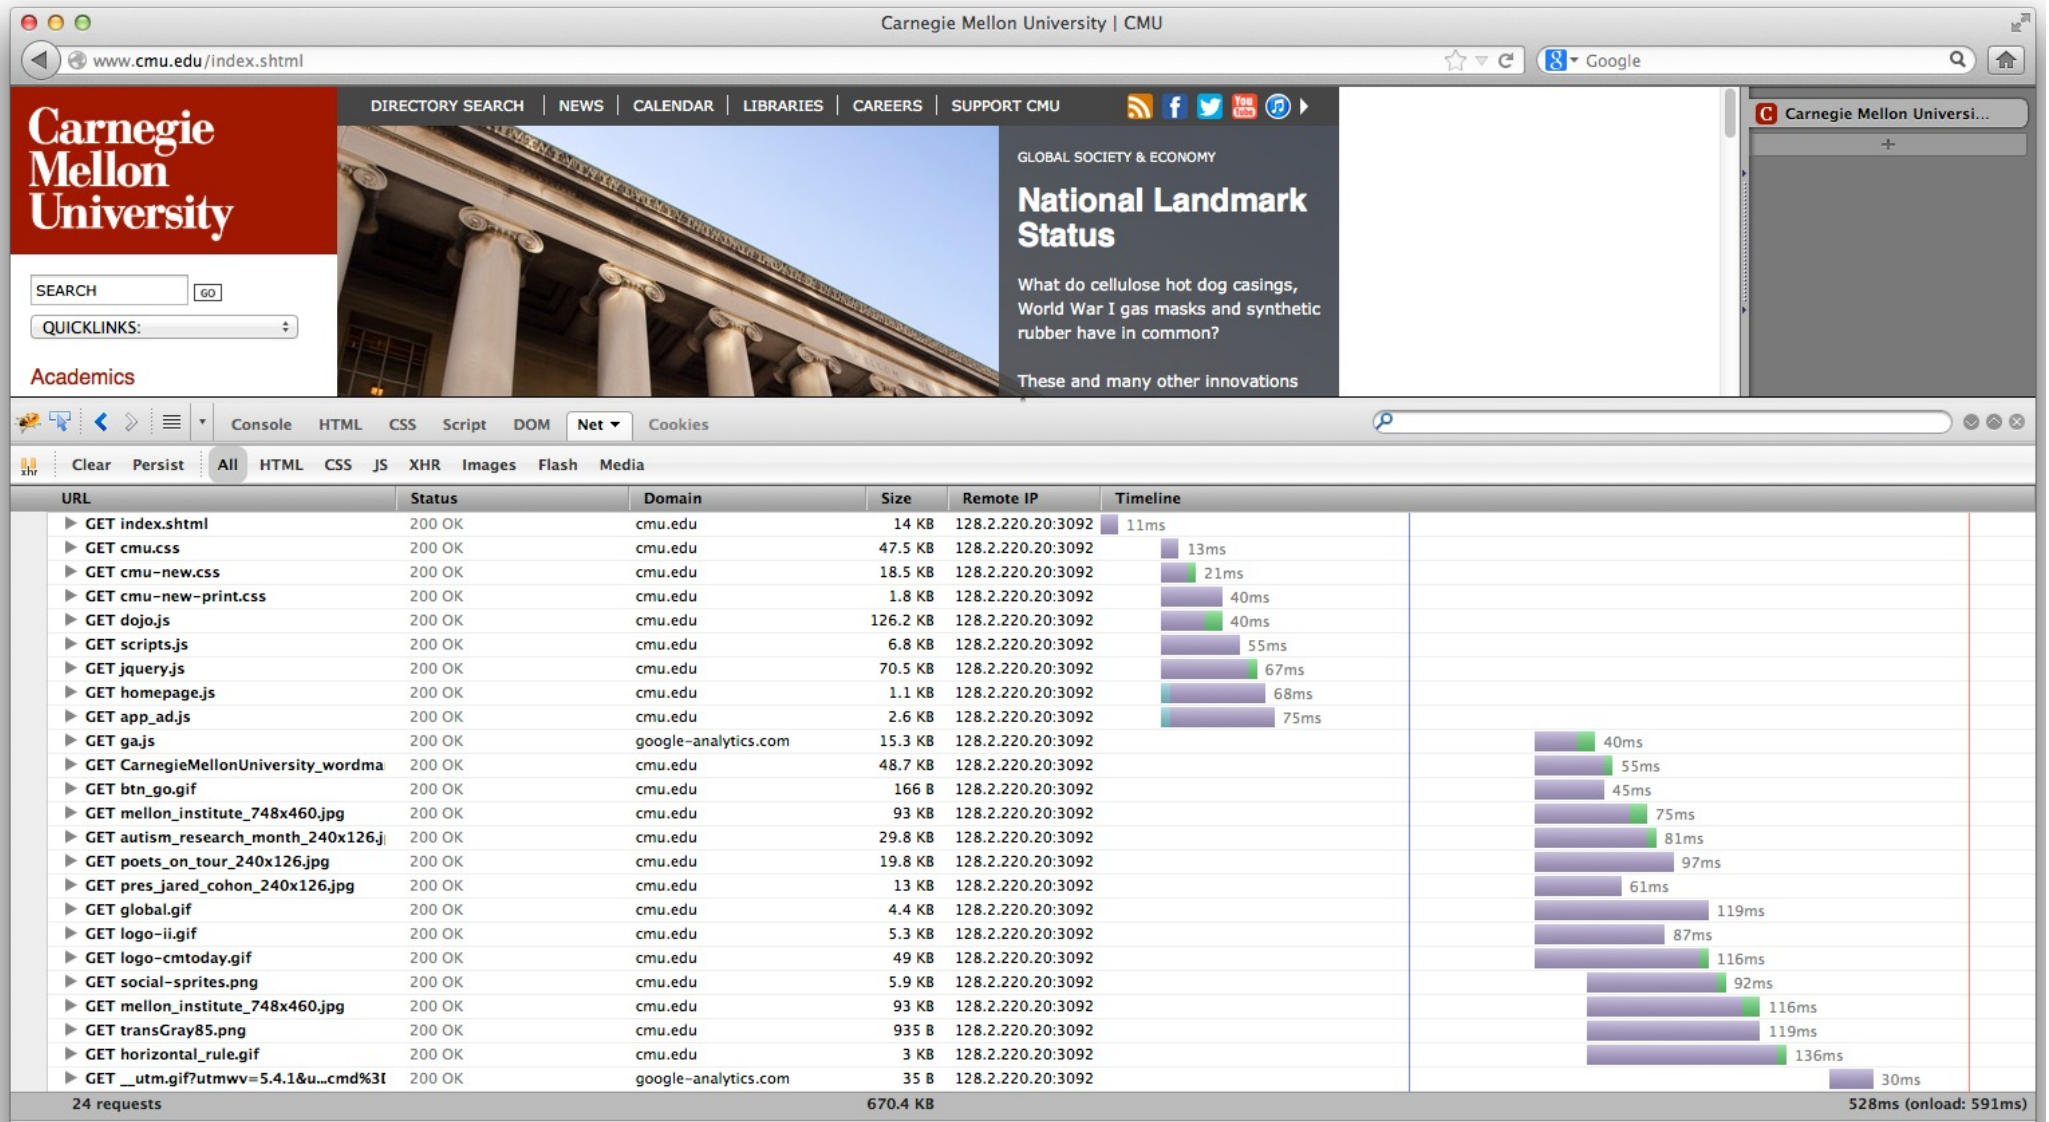
- **In reality, a single HTML page today may depend on 10s or 100s of support files (images, stylesheets, scripts, etc.)**
- **Builds a good argument for concurrent servers**
  - Just to load a single modern webpage, the client would have to wait for 10s of back-to-back request
  - I/O is likely slower than processing, so back
- **Caching is simpler if done in pieces rather than whole page**
  - If only part of the page changes, no need to fetch old parts again
  - Each object (image, stylesheet, script) already has a unique URL that can be used as a key

# How the Web Really Works

## ■ Excerpt from [www.cmu.edu/index.html](http://www.cmu.edu/index.html):

```
<html lang="en" xml:lang="en" xmlns="http://www.w3.org/1999/xhtml">
<head>
  ...
  <link href="homecss/cmu.css" rel="stylesheet" type="text/css"/>
  <link href="homecss/cmu-new.css" rel="stylesheet" type="text/css"/>
  <link href="homecss/cmu-new-print.css" media="print" rel="stylesheet"
type="text/css"/>
  <link href="http://www.cmu.edu/RSS/stories.rss" rel="alternate" title="Carnegie
Mellon Homepage Stories" type="application/rss+xml"/>
  ...
  <script language="JavaScript" src="js/dojo.js" type="text/javascript"></script>
  <script language="JavaScript" src="js/scripts.js"
type="text/javascript"></script>
  <script language="javascript" src="js/jquery.js" type="text/javascript"></script>
  <script language="javascript" src="js/homepage.js"
type="text/javascript"></script>
  <script language="javascript" src="js/app_ad.js" type="text/javascript"></script>
  ...
  <title>Carnegie Mellon University | CMU</title>
</head>
<body> ...
```

# Sequential Proxy

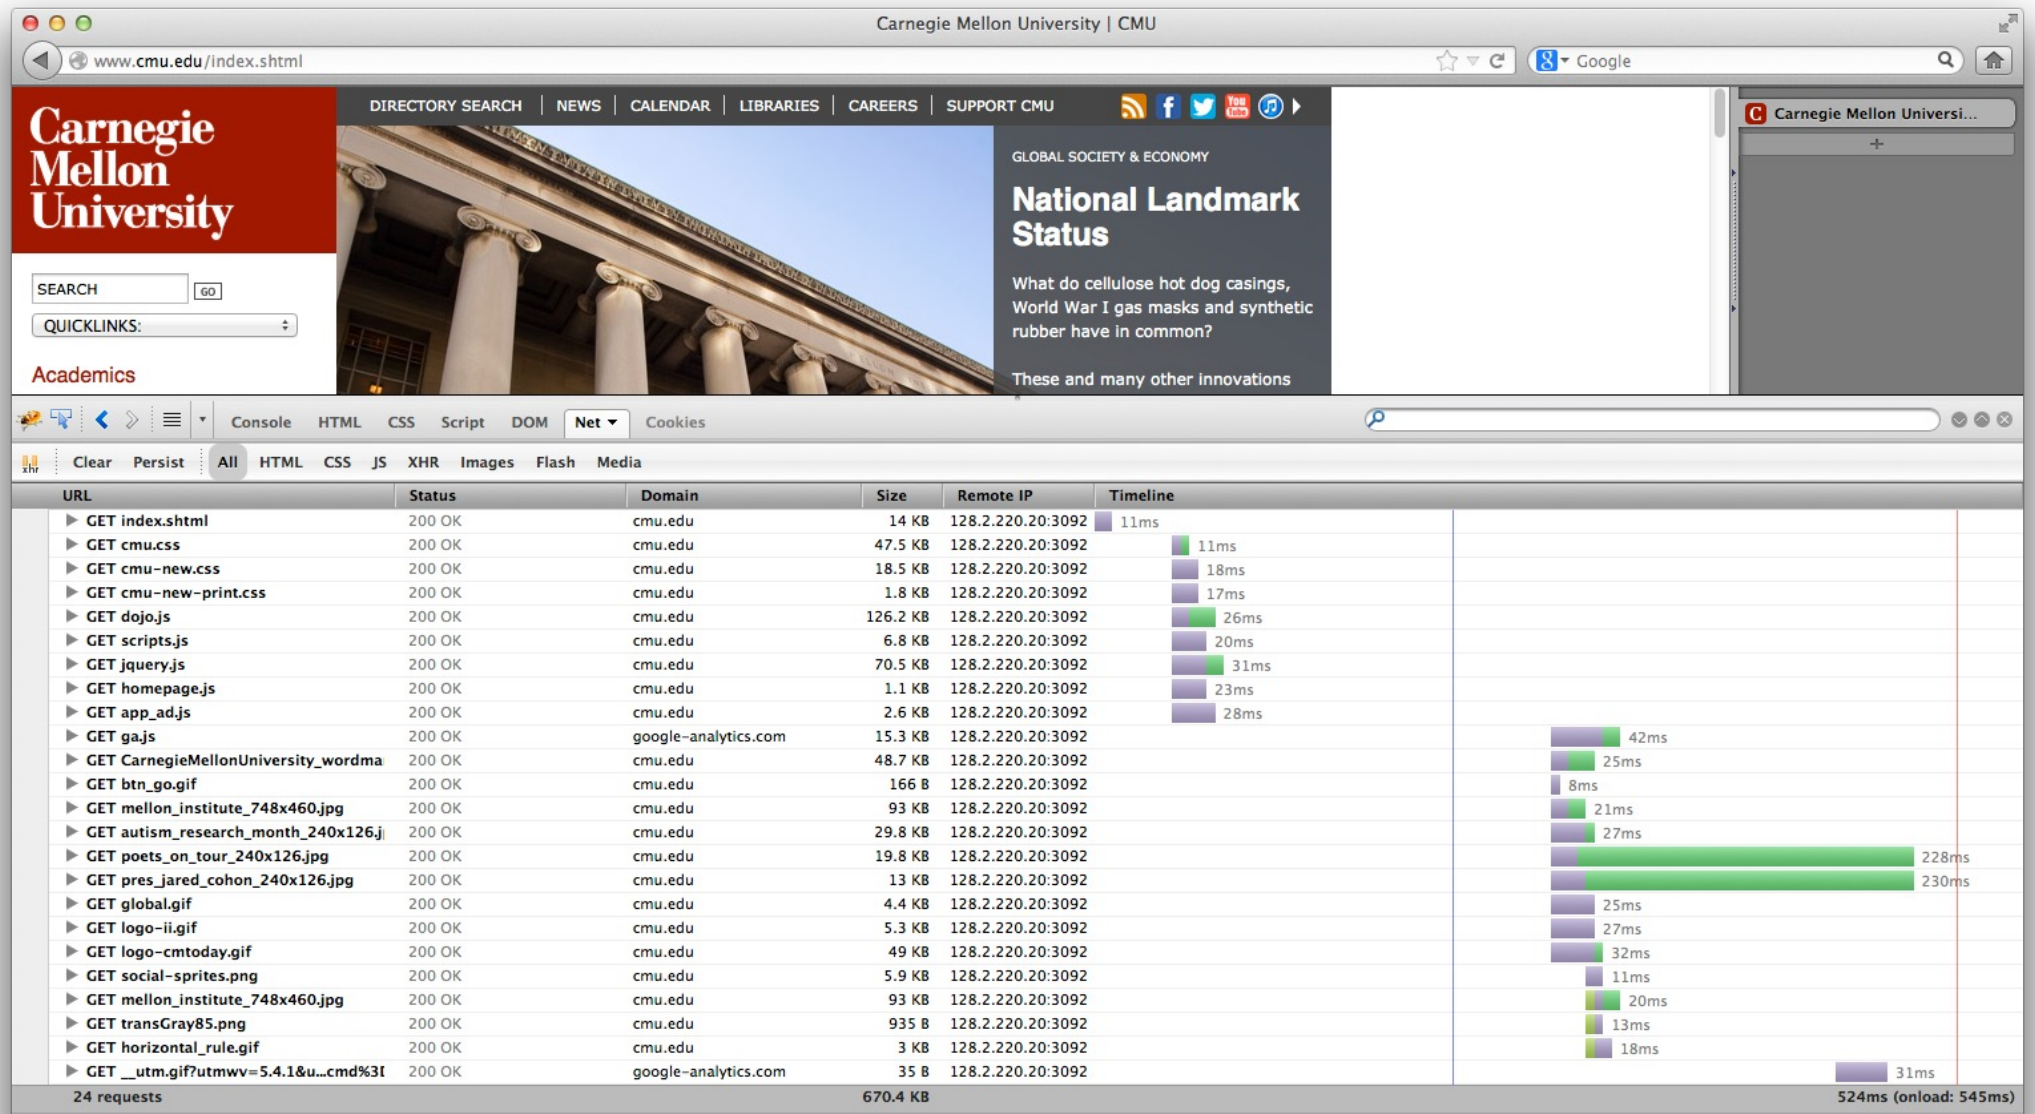




# Sequential Proxy

- **Note the sloped shape of when requests finish**
  - Although many requests are made at once, the proxy does not accept a new job until it finishes the current one
  - Requests are made in batches. This results from how HTML is structured as files that reference other files.
- **Compared to the concurrent example (next), this page takes a long time to load with just static content**

# Concurrent Proxy



# Concurrent Proxy

- Now, we see much less purple (waiting), and less time spent overall.
- Notice how multiple green (receiving) blocks overlap in time
  - Our proxy has multiple connections open to the browser to handle several tasks at once

# How the Web Really Works

## ■ A note on AJAX (and XMLHttpRequests)

- Normally, a browser will make the initial page request then request any supporting files
- And XMLHttpRequest is simply a request from the page once it has been loaded & the scripts are running
- The distinction does not matter on the server side – everything is an HTTP Request

# Outline


- **Getting content on the web: Telnet/cURL Demo**
  - How the web really works
- **Networking Basics**
- **Proxy**
  - Due Tuesday, December 8<sup>th</sup>
  - Grace days allowed
- **String Manipulation in C**

# Sockets

## ■ What is a socket?

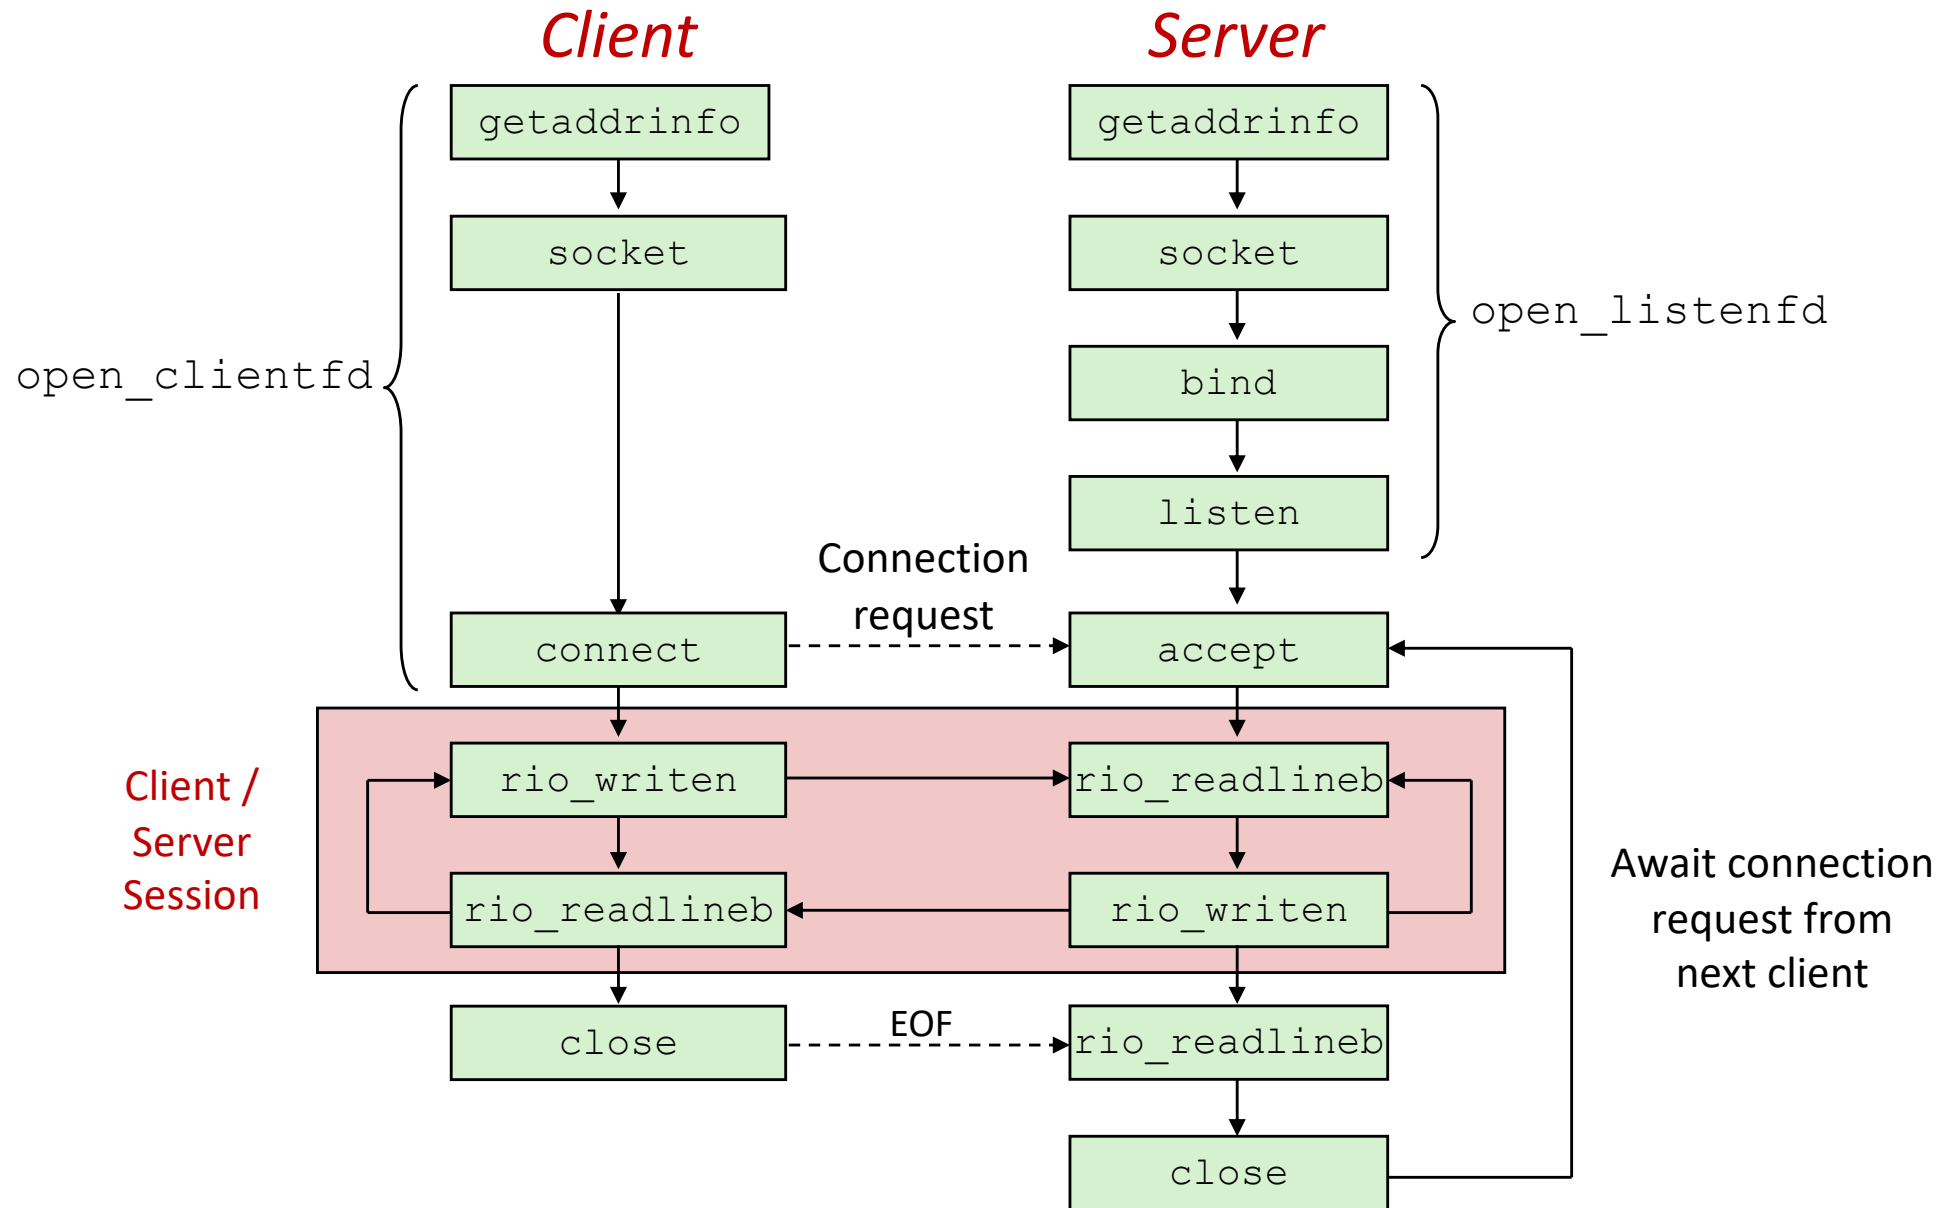
- To an application, a socket is a file descriptor that lets the application read/write from/to the network
- (all Unix I/O devices, including networks, are modeled as files)

## ■ Clients and servers communicate with each other by reading from and writing to socket descriptors



- The main difference between regular file I/O and socket I/O is how the application “opens” the socket descriptors

# Overview of the Sockets Interface



# Host and Service Conversion: `getaddrinfo`

- **`getaddrinfo`** is the modern way to convert string representations of host, ports, and service names to socket address structures.
  - Replaces obsolete `gethostbyname` - unsafe because it returns a pointer to a static variable
- **Advantages:**
  - Reentrant (can be safely used by threaded programs).
  - Allows us to write portable protocol-independent code(IPv4 and IPv6)
  - Given `host` and `service`, `getaddrinfo` returns `result` that points to a linked list of `addrinfo` structs, each pointing to socket address struct, which contains arguments for sockets APIs.
- **`getnameinfo`** is the inverse of `getaddrinfo`, converting a socket address to the corresponding host and service.



# Sockets API

- `int socket(int domain, int type, int protocol);`
  - Create a file descriptor for network communication
  - used by both clients and servers
  - `int sock_fd = socket(PF_INET, SOCK_STREAM, IPPROTO_TCP);`
  - One socket can be used for two-way communication
  
- `int bind(int socket, const struct sockaddr *address, socklen_t address_len);`
  - Associate a socket with an IP address and port number
  - used by servers
  - `struct sockaddr_in sockaddr` – family, address, port

# Sockets API

- `int listen(int socket, int backlog);`
  - `socket`: socket to listen on
  - used by servers
  - `backlog`: maximum number of waiting connections
  - `err = listen(sock_fd, MAX_WAITING_CONNECTIONS);`
  
- `int accept(int socket, struct sockaddr *address, socklen_t *address_len);`
  - used by servers
  - `socket`: socket to listen on
  - `address`: pointer to `sockaddr` struct to hold client information after `accept` returns
  - `return`: file descriptor

# Sockets API

- `int connect(int socket, struct sockaddr *address, socklen_t address_len);`
  - attempt to connect to the specified IP address and port described in address
  - used by clients
  
- `int close(int fd);`
  - used by both clients and servers
  - (also used for file I/O)
  - fd: socket fd to close

# Sockets API

- `ssize_t read(int fd, void *buf, size_t nbyte);`
  - used by both clients and servers
  - (also used for file I/O)
  - `fd`: (socket) fd to read from
  - `buf`: buffer to read into
  - `nbytes`: buf length
  
- `ssize_t write(int fd, void *buf, size_t nbyte);`
  - used by both clients and servers
  - (also used for file I/O)
  - `fd`: (socket) fd to write to
  - `buf`: buffer to write
  - `nbytes`: buf length

# Outline

- Getting content on the web: Telnet/cURL Demo
  - How the web really works
- Networking Basics
- **Proxy**
  - Due Tuesday, December 8th
  - Grace days allowed
- String Manipulation in C

# Byte Ordering Reminder

- So, how are the bytes within a multi-byte word ordered in memory?
- **Conventions**
  - Big Endian: Sun, PPC Mac, Internet
    - Least significant byte has highest address
  - Little Endian: x86, ARM processors running Android, iOS, and Windows
    - Least significant byte has lowest address

# Byte Ordering Reminder

- So, how are the bytes within a multi-byte word ordered in memory?
- Conventions
  - Big Endian: Sun, PPC Mac, **Internet**
    - Least significant byte has highest address
- **Make sure to use correct endianness**

# Proxy - Functionality

## ■ Should work on vast majority of sites

- Twitch, CNN, NY Times, etc.
- Some features of sites which require the POST operation (sending data to the website), will not work
  - Logging in to websites, sending Facebook message
- HTTPS is not expected to work
  - Google, YouTube (and some other popular websites) now try to push users to HTTPs by default; watch out for that

## ■ Cache previous requests

- Use LRU eviction policy
- Must allow for concurrent reads while maintaining consistency
- Details in write up



# Proxy - Functionality

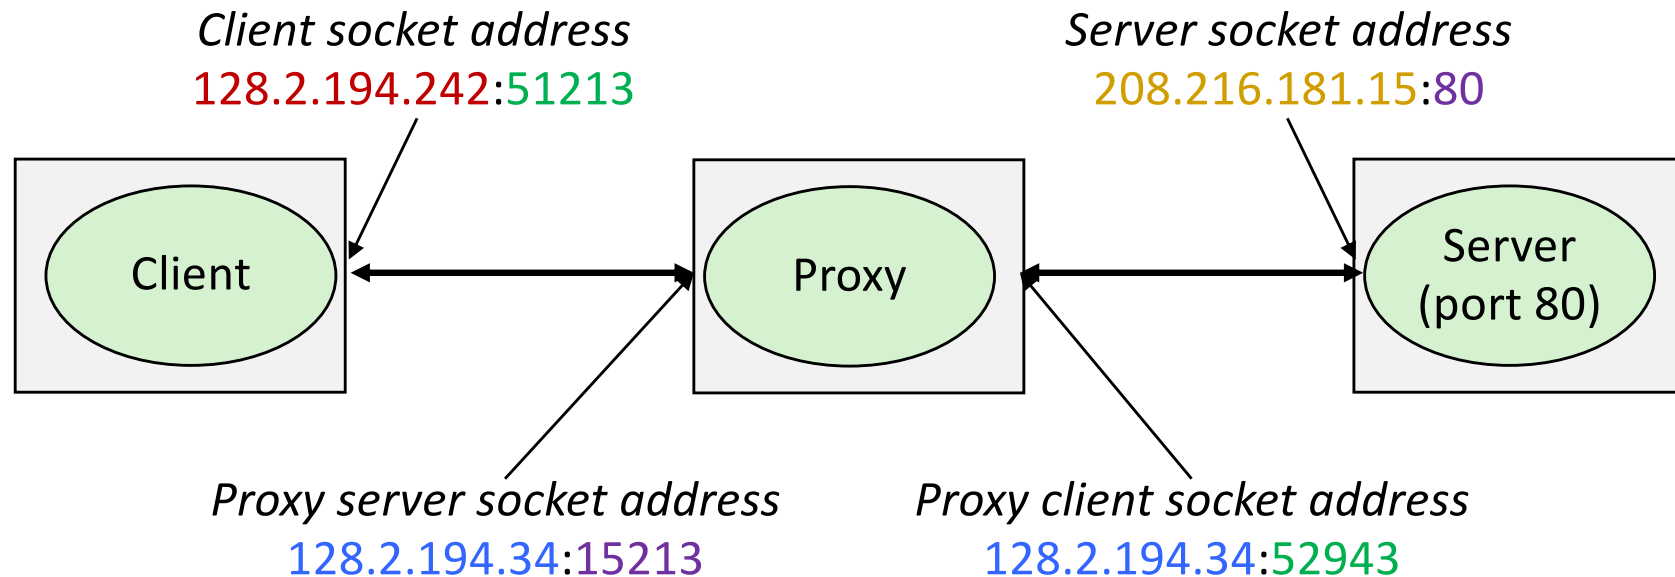
- **Why a multi-threaded cache?**
  - Sequential cache would bottleneck parallel proxy
  - Multiple threads can read cached content safely
    - Search cache for the right data and return it
    - Two threads can read from the same cache block
  - But what about writing content?
    - Overwrite block while another thread reading?
    - Two threads writing to same cache block?

# Proxy - How

- Proxies are a bit special - they are a server and a client at the same time.
- They take a request from one computer (acting as the server), and make it on their behalf (as the client).
- Ultimately, the control flow of your program will look like a server, but will have to act as a client to complete the request
- **Start small**
  - Grab yourself a copy of the echo server (pg. 946) and client (pg. 947) in the book
  - Also review the tiny.c basic web server code to see how to deal with HTTP headers
    - Note that tiny.c ignores these; you may not

# Proxy - How

- What you end up with will resemble:



# Summary

## ■ Step 1: Sequential Proxy

- Works great for simple text pages with embedded styles

## ■ Step 2: Concurrent Proxy

- multi-threading

## ■ Step 3 : Cache Web Objects

- Cache individual objects, not the whole page
- Use an LRU eviction policy
- Your caching system must allow for *concurrent reads* while maintaining consistency. Concurrency? Shared Resource?

# Outline

- Getting content on the web: Telnet/cURL Demo
  - How the web really works
- Networking Basics
- Proxy
  - Due Tuesday, December 8th
  - Grace days allowed
- **String Manipulation in C**

# String manipulation in C

## ■ sscanf: Read input in specific format

```
int sscanf(const char *str, const char *format, ...);
```

Example:

```
buf = "213 is awesome"
```

```
// Read integer and string separated by white space from buffer 'buf'
```

```
// into passed variables
```

```
ret = sscanf(buf, "%d %s %s", &course, str1, str2);
```

This results in:

```
course = 213, str1 = is, str2 = awesome, ret = 3
```

# String manipulation (cont)

- **sprintf: Write input into buffer in specific format**

```
int sprintf(char *str, const char *format, ...);
```

Example:

```
buf[100];
```

```
str = "213 is awesome"
```

```
// Build the string in double quotes ("" ) using the passed arguments
```

```
// and write to buffer 'buf'
```

```
sprintf(buf, "String (%s) is of length %d", str, strlen(str));
```

This results in:

```
buf = String (213 is awesome) is of length 14
```

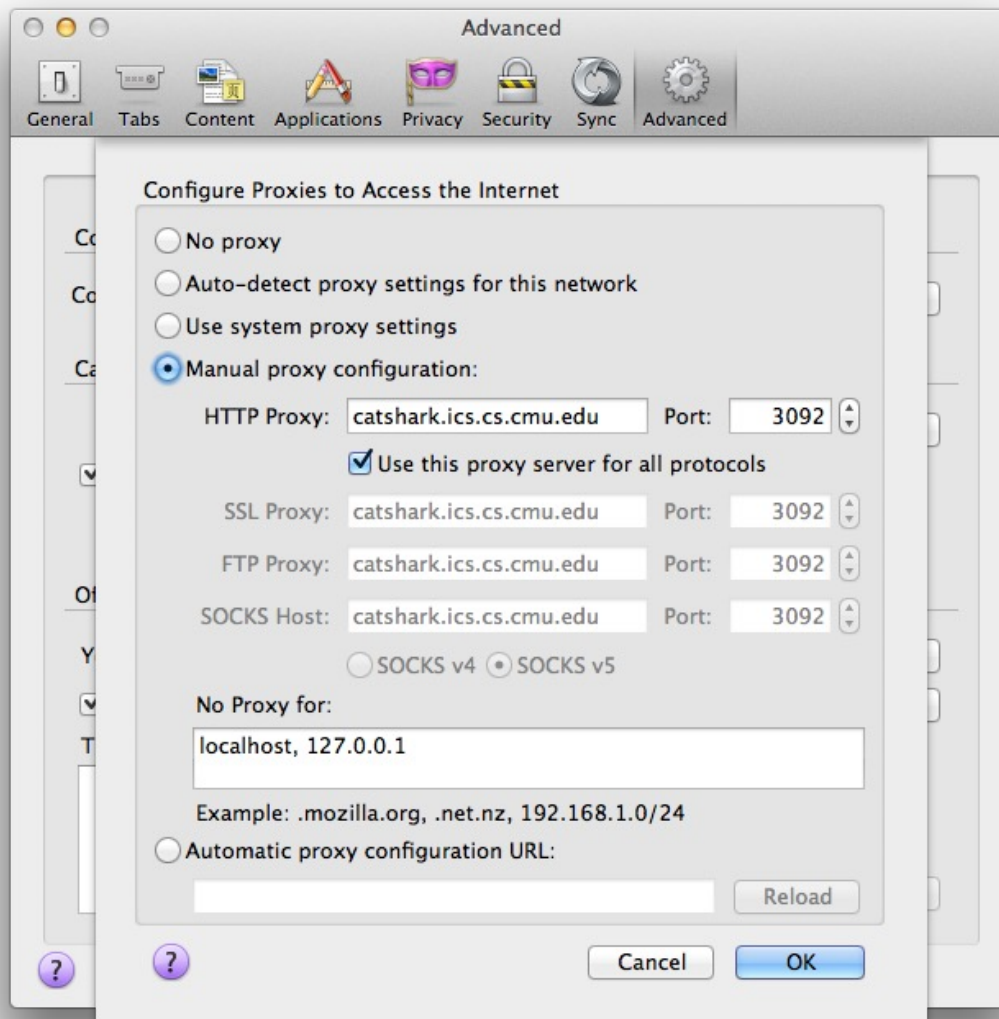
# String manipulation (cont)

Other useful string manipulation functions:

- `strcmp`, `strncmp`, `strncasecmp`
- `strstr`
- `strlen`
- `strcpy`, `strncpy`



# Aside: Setting up Firefox to use a proxy



- You may use any browser, but we'll be grading with Firefox
- Preferences > Advanced > Network > Settings... (under Connection)
- Check "Use this proxy for all protocols" or your proxy will appear to work for HTTPS traffic.